

Model Comparison

VLM vs external model on the binary task

Discrete metrics at the 0.5 threshold

Metric	VLM	External	Delta
accuracy	0.9005	0.9285	-0.0280
macro_f1	0.8999	0.9281	-0.0282
weighted_f1	0.8999	0.9281	-0.0282

Per-sample summary

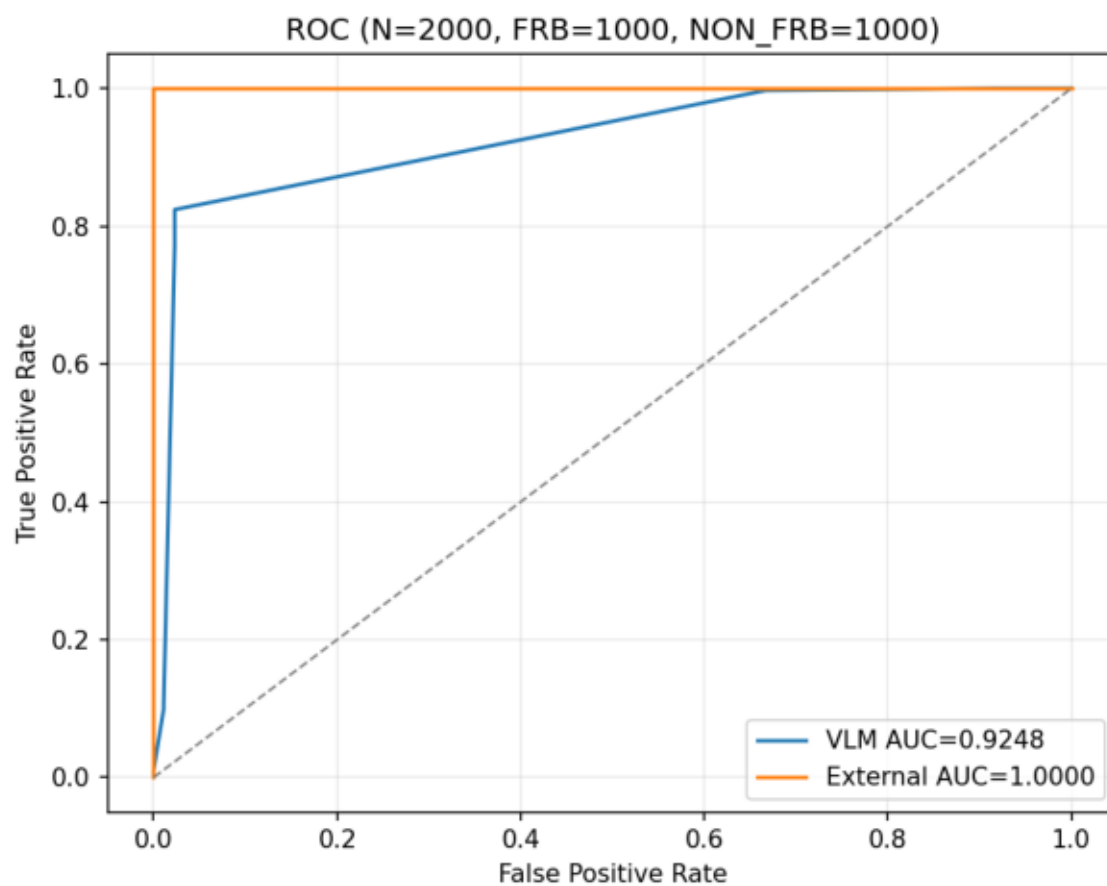
Indicator	Value
Common samples	2000
Agreements	1700
Disagreements	300
Only the VLM correct	122
Only the external model correct	178
Both correct	1679
Both wrong	21
Samples only in the VLM file	0
Samples only in the external file	0

Binary probabilistic analysis

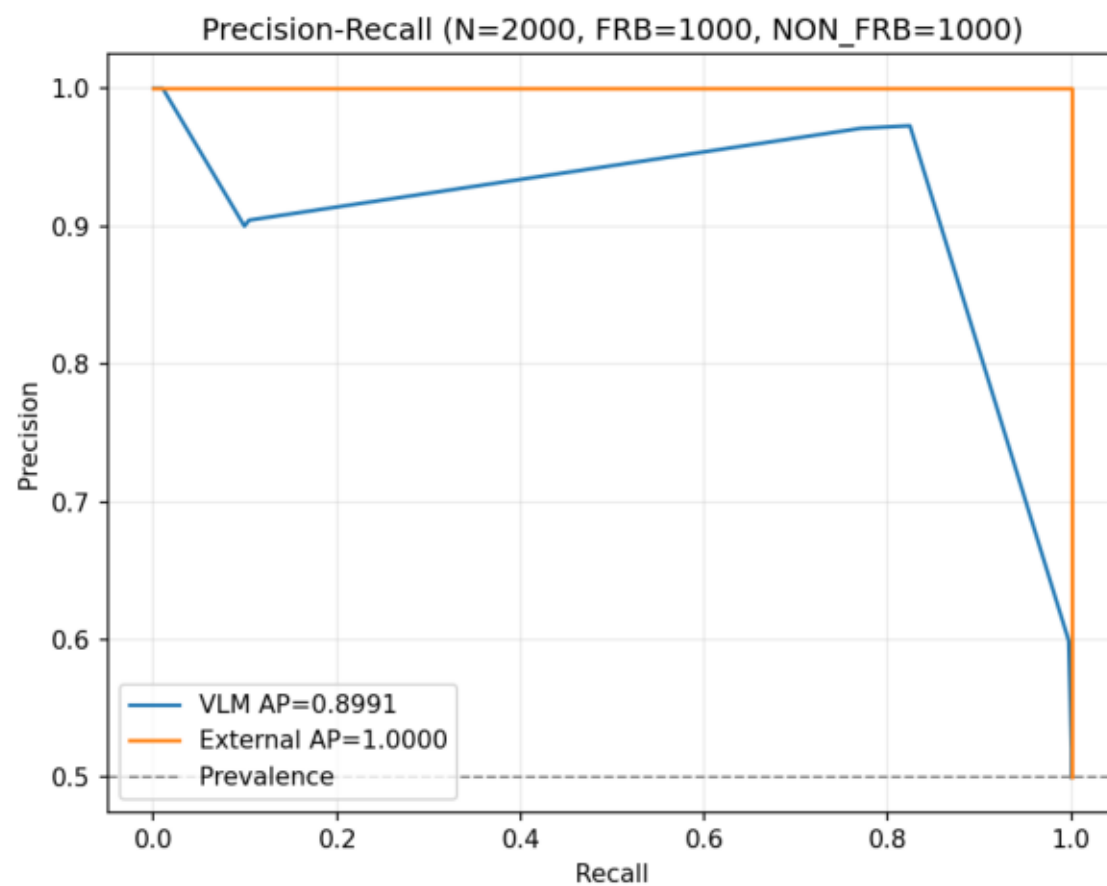
Indicator	Value
Paired samples	2000
FRB	1000
NON_FRB	1000
AP VLM	0.8991 [0.8814, 0.9151]
AP external	1.0000 [1.0000, 1.0000]
Delta AP	-0.1009 [-0.1186, -0.0849]
AUC VLM	0.9248 [0.9141, 0.9351]
AUC external	1.0000 [1.0000, 1.0000]

The ROC/PR curves use `parsed_response.frb_probability`. The root-level confidence is only the confidence in the discrete label. Log-loss, Brier, and ECE use clipping to [1e-7, 1-1e-7].


ROC Curve




Precision-Recall Curve



Score Histogram



Operating Point



Calibration

