

Evaluation Report

Prediction summary

Summary

Indicator	Value
Total samples	3000
Evaluated samples	3000
Skipped samples	0
Model errors	0
Invalid labels	0
Accuracy	0.8013
Macro-F1	0.7987
Weighted-F1	0.7987

Per-class metrics

Class	Precision	Recall	F1-score	Support
FRB	1.0000	0.8700	0.9305	1000
RFI	0.9889	0.5340	0.6935	1000
NOISE	0.6289	1.0000	0.7722	1000

Performance interpretation

The model showed good performance, with accuracy 0.8013 and macro-F1 0.7987 over the 3000 evaluated samples. The best class by F1-score was FRB (0.9305), while the weakest class was RFI (0.6935). Macro-F1 and weighted-F1 are close, suggesting relatively balanced performance across the evaluated classes. The most frequent confusion was RFI classified as NOISE, with 466 occurrence(s).

How to read the metrics

- Accuracy: fraction of evaluated samples whose predicted class matches the true label.
- Per-class precision: among all images predicted as a class, measures how many actually belonged to that class.
- Per-class recall: among all images that actually belonged to a class, measures how many the model recovered.
- Per-class F1-score: harmonic mean of precision and recall; penalizes scenarios where either metric is low.
- Macro-F1: simple mean of the F1-scores of the classes FRB, RFI, NOISE; every class has the same weight.
- Weighted-F1: mean of the F1-scores weighted by the number of samples of each class.
- Support: number of evaluated samples of each true class.
- Confusion matrix: rows represent the true class and columns represent the predicted class.

Confusion Matrix

