

Evaluation Report

Prediction summary

Summary

Indicator	Value
Total samples	2000
Evaluated samples	2000
Skipped samples	0
Model errors	0
Invalid labels	0
Accuracy	0.9005
Macro-F1	0.8999
Weighted-F1	0.8999

Per-class metrics

Class	Precision	Recall	F1-score	Support
FRB	0.9728	0.8240	0.8923	1000
NON_FRB	0.8474	0.9770	0.9076	1000

Performance interpretation

The model showed high performance, with accuracy 0.9005 and macro-F1 0.8999 over the 2000 evaluated samples. The best class by F1-score was NON_FRB (0.9076), while the weakest class was FRB (0.8923). Macro-F1 and weighted-F1 are close, suggesting relatively balanced performance across the evaluated classes. The most frequent confusion was FRB classified as NON_FRB, with 176 occurrence(s).

How to read the metrics

- Accuracy: fraction of evaluated samples whose predicted class matches the true label.
- Per-class precision: among all images predicted as a class, measures how many actually belonged to that class.
- Per-class recall: among all images that actually belonged to a class, measures how many the model recovered.
- Per-class F1-score: harmonic mean of precision and recall; penalizes scenarios where either metric is low.
- Macro-F1: simple mean of the F1-scores of the classes FRB, NON_FRB; every class has the same weight.
- Weighted-F1: mean of the F1-scores weighted by the number of samples of each class.
- Support: number of evaluated samples of each true class.
- Confusion matrix: rows represent the true class and columns represent the predicted class.

Confusion Matrix

